

JANHAVI SHAILENDRA JADHAV

Baltimore, MD | (443) 418-0917 | janhavij0603@gmail.com | linkedin.com/in/janhavi-jadhav006

PROFESSIONAL SUMMARY

Graduate student who has shipped 10+ production integrations processing 5M+ records/day - connecting enterprise SaaS systems via REST APIs, triaging schema mismatches and data variances in real time, and publishing validated, application-ready outputs to downstream consumers. Built reusable transformation recipes across Snowflake, dbt, PySpark, and Airflow; delivered integration flows end-to-end from source mapping through stakeholder sign-off. Comfortable owning integrations independently, translating messy real-world data into trusted models, and closing the loop with both technical and non-technical teams.

TECHNICAL SKILLS

- **Languages & Query:** Python, SQL (Advanced), TypeScript, Java
- **Data Integration & APIs:** REST API design & consumption, sFTP, API-based data connections, JSON/CSV/TSV/Excel file handling
- **Data Warehouses:** Snowflake, BigQuery, PostgreSQL, MySQL, SQLite
- **Transformation & Modeling:** PySpark, dbt, Pandas, data cleansing, normalization, schema mapping, semi-structured data
- **Orchestration & Streaming:** Apache Airflow, Apache Kafka, CI/CD, Docker
- **Frameworks & APIs:** FastAPI, Flask, Streamlit, LangChain, RAG Pipelines
- **Spreadsheet & Validation:** Advanced Microsoft Excel, data validation, audit matching, reconciliation
- **AI / LLM Exposure:** Anthropic API, OpenAI API, Groq API, automated data prep, schema mapping with LLMs
- **Collaboration & Delivery:** Cross-functional stakeholder management, technical documentation, Agile, post-sales delivery mindset

PROFESSIONAL EXPERIENCE

Tata Consultancy Services

January 2024 – July 2024

Assistant Systems Engineer | Data Integration & Pipeline Engineering

- Reduced downstream data errors by 30% across 5M+ records/day by engineering reusable end-to-end integration pipelines - extracting from upstream enterprise SaaS systems, applying standardized transformation and cleansing recipes, and publishing validated outputs to ML and reporting applications.
- Achieved complete data trust across 10+ production integrations by establishing REST API-based connections to enterprise source systems, triaging schema mismatches and data variances end-to-end, and delivering validated, application-ready data flows to downstream stakeholders.
- Onboarded new source systems 35% faster by designing standardized, reusable transformation recipes - with automated audit-matching and validation gates that caught schema gaps, normalization errors, and incomplete source records before they reached downstream consumers.
- Closed the product feedback loop by surfacing front-line integration pain points - schema edge cases, mapping ambiguities, and source-system quirks - to Engineering and Product, directly influencing two platform-level improvements that reduced manual triage time across the team.
- Accelerated integration launch timelines by partnering cross-functionally with Engineering, Product, and customer-facing delivery teams to operationalize data flows, writing technical documentation and mapping guides for each REST API integration deployed to production.

TCR Innovation

March 2022 – May 2023

Data & Analytics Engineering Intern

- Delivered clean, model-ready datasets across 1M+ records by designing and executing analytics workflows covering data cleansing, normalization, schema mapping, and reconciliation - reducing downstream rework by standardizing transformation logic across sources.
- Extracted structured insights from semi-structured and unstructured sources by applying NLP techniques (tokenization, embeddings, text classification), with findings communicated directly to non-technical business stakeholders to inform data-driven decisions.
- Ensured 90%+ precision and recall on model outputs by building validation pipelines with audit-matching metrics (precision, recall, F1), preparing Excel-based reconciliation reports for stakeholder review, and documenting all findings for reproducibility.

- Surfaced analytics results to business teams by developing dashboards that translated complex model outputs into actionable metrics, improving visibility into data quality and supporting decision-making for both technical and non-technical stakeholders.

MicDots

January 2026 – May 2026

Project Management Intern

- Built data-backed reporting dashboards and KPI tracking systems for technology pilot rollouts, turning scattered project metrics into a single source of truth used by both engineering and business stakeholders to monitor delivery progress.
- Served as the technical-to-business translation layer during product implementations - gathering requirements from non-technical stakeholders, scoping deliverables with engineering, and tracking open issues through to resolution to keep launches on schedule.

PROJECTS

Market Intelligence Data Platform - End-to-End Data Integration & Analytics | [GitHub](#) June 2026

- Built a full streaming financial data platform using Kafka for real-time ingestion, PySpark for distributed transformation, and Snowflake as the cloud data warehouse - mirroring the end-to-end integration flow from source to application-ready model.
- Modeled analytics-ready datasets with dbt (normalization, reshaping, business logic) and exposed them through a FastAPI service and Apache Superset dashboards for investment signal analysis.
- Triaged and resolved complex infrastructure issues including Kafka listener misconfigurations, Snowflake MFA bypass for service accounts, and Airflow executor compatibility - demonstrating a strong production troubleshooting mindset.
- Implemented Airflow DAGs for orchestration and CI/CD pipelines via GitHub Actions for reliable, repeatable data delivery.

HealthFlow - Cloud-Native Healthcare Claims Analytics Platform | [GitHub](#)

June 2026

- Architected a cloud-native data platform on AWS and GCP using Terraform for infrastructure provisioning, PySpark for claims data transformation, and dbt on BigQuery for analytics modeling.
- Built a FastAPI serving layer exposing validated, analytics-ready claims data to downstream applications, with end-to-end pipeline orchestration via Apache Airflow on Kubernetes.
- Designed data ingestion workflows handling structured and semi-structured healthcare claims data with validation, schema normalization, and quality checks at each transformation stage.

FinSight - Multi-Agent Financial Research System | [GitHub](#)

May 2026

- Architected a multi-agent system orchestrating 4 specialized agents to generate cited financial analyst reports from SEC 10-K filings (structured + semi-structured source data) in under 12 seconds with streaming delivery via FastAPI.
- Implemented hybrid RAG combining ChromaDB vector search across 1,138 chunks with a NetworkX knowledge graph of 366 risk nodes - demonstrating AI/LLM tooling applied to automated data extraction and schema mapping from complex financial documents.
- Engineered an automated evaluation pipeline with prompt injection guardrails, full SQLite audit logging, and 90%+ test coverage; exposed via structured-output FastAPI REST API.

Product Sentiment Analysis Pipeline | [GitHub](#)

January 2026

- Developed a data pipeline processing 500K+ customer reviews end-to-end: extraction, cleansing, transformation, RoBERTa-based classification, and publishing to a PostgreSQL store exposed via REST API.
- Containerized all services with Docker and implemented CI/CD pipelines for reliable delivery; designed monitoring dashboards for operational visibility and data quality metrics.

EDUCATION

University of Maryland, Baltimore County

May 2026

Master of Science in Information Systems | GPA: 3.97/4.0

University of Mumbai

May 2023

Bachelor of Engineering in Electronics | GPA: 3.42/4.0